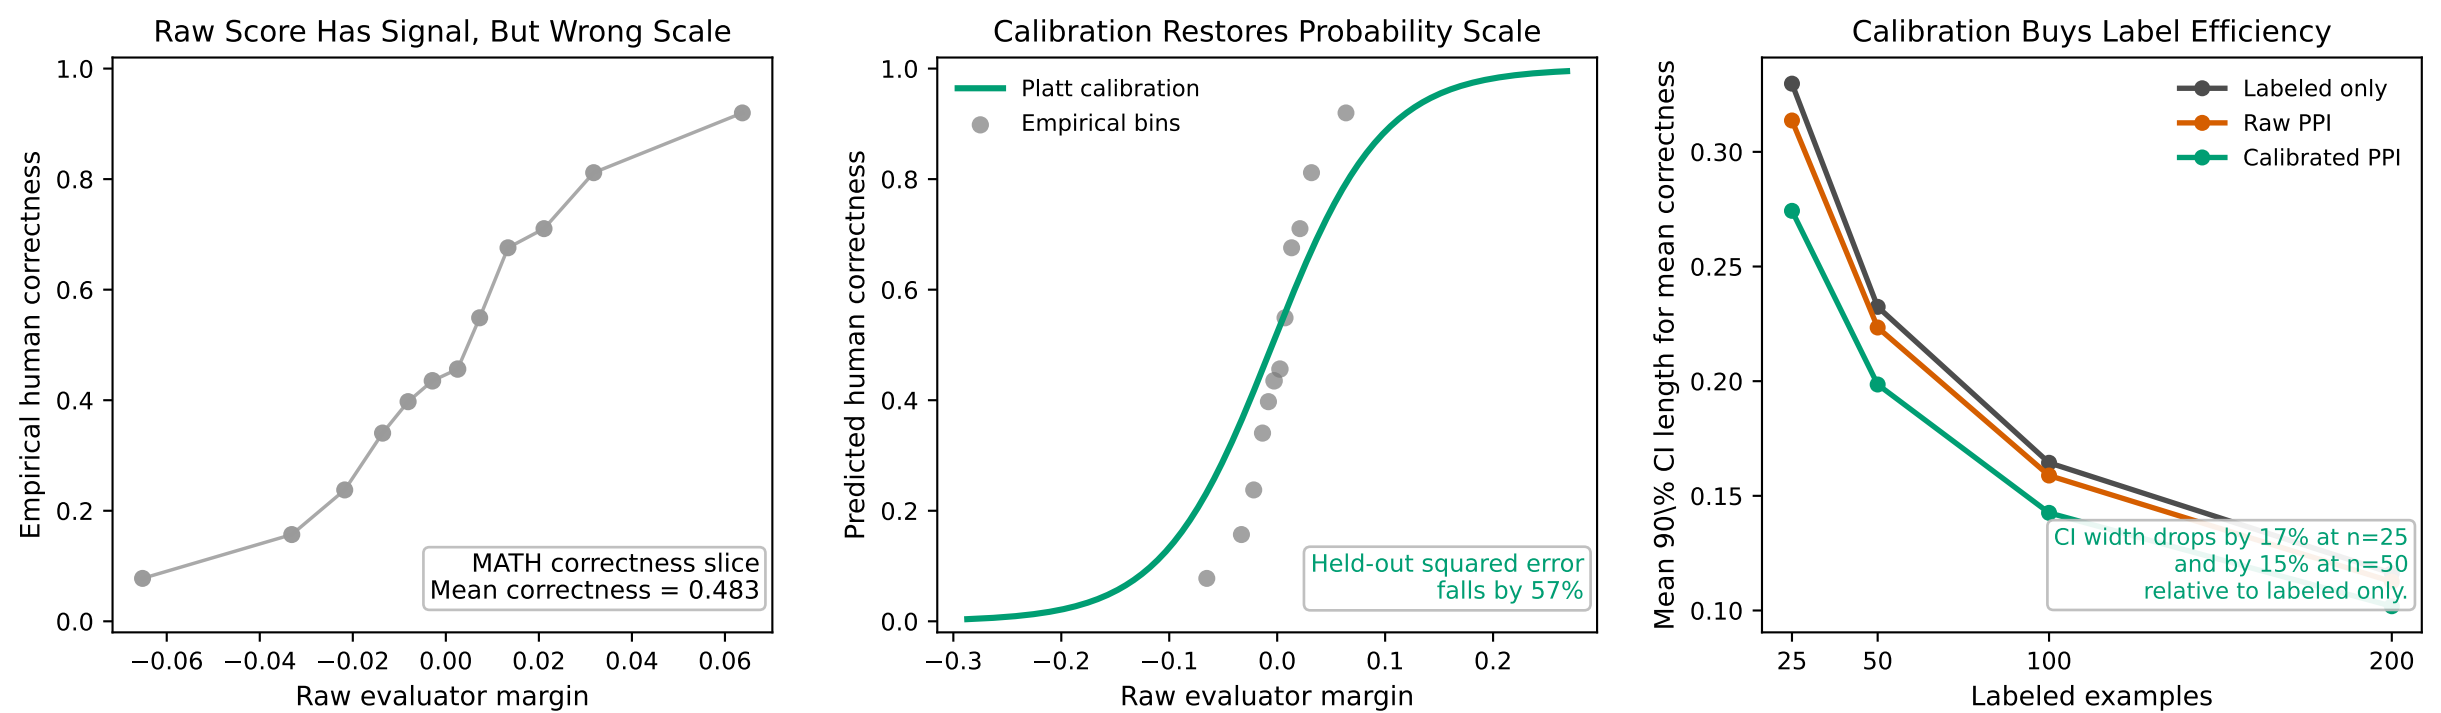
Raw Score Has Signal, But Wrong Scale
Empirical human correctness
Raw evaluator margin
MATH correctness slice
Mean correctness = 0.483
Calibration Restores Probability Scale
Predicted human correctness
Raw evaluator margin
Platt calibration
Empirical bins
Held-out squared error
falls by 57%
Calibration Buys Label Efficiency
Mean 90\% CI length for mean correctness
Labeled examples
Labeled only
Raw PPI
Calibrated PPI
CI width drops by 17% at n=25
and by 15% at n=50
relative to labeled only.